Predicting the outcome of a PCR test for COVID-19 using a routine blood exam: A replication

An example of how AI can help identify possible infections in a primary care/triage system

By

Lukas Schießer

Abstract

Acknowledgements

I want to thank...

Contents

List of Figures

List of Tables

Chapter 1

Introduction

Introduction for this thesis by explaining motivation for this topic and giving a short overview.

Motivation: overwhelmed health care systems, shortages of testing supplies, AI as supplementary technique to help.

Chapter 2

Materials and Methods

2.1 Data

The data used to train the classifiers was provided by Brinati et al. [5]. It was collected between the end of February 2020 and mid of March 2020 from patients admitted to the *IRCSS Ospedale San Raffaele* and consists of 279 individuals who were selected randomly. For each individual, the data set provides their age, gender, results of a routine blood screening, and the result of a PCR test for Sars-CoV-2. A complete overview over the recorded variables is provided in A.1. The target variable *Swab* is binary and indicates the result of a PCR-test for Sars-CoV-2 taken by naso-pharyngeal swab. A 0 indicates a negative test while a 1 indicates a positive test. The data set is slightly imbalanced towards positive cases with 102 (37%) negative cases and 177 (63%) positive cases.

Since the variable *Gender* was provided as a string, it was transformed into two binary numerical variables called *female* and *male* by one-hot encoding. Further, two values of the variable *Age* were removed, specifically the values 0 and 1. This was sensible seeing that there was no other data recorded from minors under the age of 18 and thus these two values can be presumed to be input errors during the collection process. Further, the variable *Lymphocytes* had to be recast using `pandas.to_numeric` because there was an input error for one of its values rendering the contents of the column a string instead of a numerical value. This process created one missing values in place of the erroneous value.

Table A.4 provides common statistics for the numerical features of the data set. All features are skewed. Blood values are skewed positively exhibiting a left-leaning distribution because for most blood values a lower or centered value is significant for a better health. Age is skewed negatively exhibiting a right-leaning distribution since COVID-19 affected older individuals more severely than younger individuals especially in Italy. As you can see in Figure B.1, most of the data is non-normally

distributed. By inspecting the kernel density plots provided in Figure B.1, it is apparent that most of the features are non-normally distributed. Additionally, a Shapiro-Wilk test for normality was performed to confirm this observation, its results are provided in table A.5. The test tests the null-hypothesis that the samples of a variable come from a normally distributed population. If the p-value of the test is smaller than 0.05 which is a commonly chosen alpha level the null-hypothesis is rejected meaning the data is not normally distributed.

Table A.3 shows that most features have missing values. 196 samples have at least one feature missing which amounts to 70 % of the data. Due to the small size of data set it is not feasible to exclude these individuals from the analysis process. It is rather more constructive to use an imputation method that models the missing values based on the observed values in the data set. Therefore, Brinati et al. chose to use *Multivariate Imputation by Chained Equations*.

2.2 Multivariate Imputation by Chained Equations

Multivariate Imputation by Chained Equations or *MICE* for short is an imputation method proposed by Buuren and Groothuis-Oudshoorn [7], it is also known as fully conditional specification (FCS). MICE is a method that imputes missing data by estimating a set of possible values from distributions of observed data. Each variable with missing data x_n is regressed on all other variables $x_1, ..., x_k$ which are restricted to the occurrences with observed data in x_n.

The imputation process is based on the following four main steps [13, 3]: Firstly, all missing values are imputed using a simple imputation method (e.g. mean imputation). These imputations can be thought of as "place holders" used during the first modeling phase. During step 2, the "place holder" imputations for one variable x are set back to missing. In step 3, all observed values from variable x in step 2 are regressed on the other variables in the imputation model. Since this is the model building phase, this step only uses samples where x has observed values. Therefore, x is the dependent variable and all other variables are independent variables used in the regression model. In step 4, the missing values in x are replaced with imputations (predictions) from the regression model built in step 3. All values of x, the observed and the imputed values, are then used in subsequent regression models of other variables.

Steps 2-4 are repeated for every variable with missing data. After the algorithm is done cycling through all variables, one iteration or "cycle" is completed. Steps 2-4 are repeated for a user-specified number of cycles. Generally, ten to twenty cycles

should suffice to stabilize the results of the imputation that is the parameters controlling the imputations should have converged by then. This imputation process is usually repeated m times creating m slightly differently imputed data sets which are then used in the subsequent analysis. According to [6, 3, 13], already a small number of imputed data sets, usually three to ten, is sufficient to provide sensible results during analysis. It can, of course, be advantageous to use a higher number of imputed data sets to get a broader range of estimates, but setting m higher requires more computations and storage and may not be worth preoccupying these extra resources [6].

MICE assumes that the data is missing at random (MAR) that is the probability of data being missing does not depend on the unobserved data but is only dependent (conditional) on the observed data. Due to its individualistic approach, MICE can handle variables of different types using different modeling choices for different variable types.

> Describe MICE algorithm (keine Bewertung, findet in Discussion statt)
> Short description of implementation in Python using rpy2
> Shortly mention that R implementation is not able to apply MICE model to other data only to data it is "trained"/ "fitted" on (or maybe that's for the 4 Discussion section)
> Maybe include PMM? probably not since this is too deep

2.3   Classifiers

The original paper [5] implements 7 classifiers, namely Random Forest, logistic regression, decision tree, k-nearest neighbors, naive Bayes, support vector machines, extremely randomized tree, and three-way random forest. Here, I will reproduce Random Forest, logistic regression, and decision tree. The following section will introduce them.

2.3.1   Decision Tree

In order to motivate the use of Random Forests and since we will use a Decision Tree during the validation to gain some insights into the workings of the Random Forest classifier, we will take a closer look at Decision Trees since they are the basic building block of Random Forests. A Decision Tree has a flowchart-like structure where each internal or decision node tests an attribute, each branch corresponds to one attribute value and each leaf node represents a classification. It is built up using the ID3 algorithm.

ID3 tries to determine the best attribute of a given data set by the distribution of its values. The best attribute is then used as a root of the decision tree and a branch is created for every value this attribute can take which also creates a subset of the data set that only has the attribute value of the branch. This process is repeated for every branch with the remaining subset of the training set until a leaf node is reached.

ID3 can use many different measures to decide which is the best attribute two of the most popular ones are information gain or the Gini Impurity.

Information gain uses the entropy measure to compute the impurity in an attribute. Entropy originates from information theory and describes the average information content of an attribute's possible outcomes, it is calculated as shown in Equation 2.1. A high entropy represents a high average information content in a attribute. Information Gain measures the expected reduction in entropy of a set S caused by learning the state of a random variable A. It is calculated as shown in Equation 2.2 and can be described as the difference between the entropy of a set S and the weighted average of the child entropies. This is computed for all remaining attributes, the attribute that maximizes the differences is then selected as a new node.

$$E(S) = \sum_{i=1}^{C} -p(i) \ log_2(p(i)) \tag{2.1}$$

$$Gain(S, A) = E(S) - \sum_{i}^{C} \frac{|S_i|}{S} E(S_i) \tag{2.2}$$

Gini Impurity or Gini index measures the impurity in a set S that is the probability of an incorrect classfication of random sample. It is calculated as shown in Equation 2.3. The measure which the algorithm bases its decision on which attribute to select as a node is called the Gini gain. It is calculated as shown in Equation 2.4 and can be described as the average Gini impurity. For the Gini gain the attribute with lowest value is selected since we want to minimize the incorrect classification of attributes.

$$G = 1 - \sum_{i=1}^{C} p(i)^2 \tag{2.3}$$

$$Gini(S, A) = \sum_{i=1}^{C} \frac{|S_i|}{|S|} G(S_i) \tag{2.4}$$

ID3 is not able to process numerical attributes and missing values. It must therefore be extended to be able to deal with real-world data by the C4.5 algorithm.

2.3.2 Random Forest

The Random Forest Classifier is an ensemble classifier that uses multiple Decision Tree instances to classify the given data. It uses two methods to diversify the different classification of the DT instances called tree bagging and feature bagging.[4] During tree bagging, the decision tree is not trained on the whole training set, but on a subset of the training set. The subsets are generated by sampling the original data set with replacement where a certain percentage of samples is selected from the original data set and the remaining percentage are duplicates of the already selected samples. Thereby this process is generating new bagged training sets S_i with the same size as the original training set S. If the size of the bagged set is smaller than the size of the original set, then this process is called sub-bagging.[11]

Feature bagging or feature subset selection limits the number of features the individual decision trees are able to use at each new split in consideration. For each decision node consideration, a number of features are randomly selected and presented to the decision tree algorithm. The number of features available to the algorithm is controlled by a hyperparameter. For classification problems, this is usually set to $\sqrt{p}$ where p is the number of features present in the complete data set.

Each decision tree in the forest generates a classification, the final classification of the Random Forest is determined by a majority vote over all decision trees. By using tree and feature bagging Random Forests average over all individual decision tree models and thereby reducing the variance of classification and trying to avoid overfitting.[2]

2.3.3 Logistic regression

Logistic regression, especially binary logistic regression, is a modification of linear regression that models the probability $p(X)$ that a sample X has the label 1. To do so, it uses a special form of the sigmoid function called the logistic function:

$$p(\mathbf{X}) = \phi_{sig}(\mathbf{X}) = \frac{e^{\mathbf{X}\vec{\beta}}}{1 + e^{\mathbf{X}\vec{\beta}}} = \frac{1}{1 + e^{-\mathbf{X}\vec{\beta}}} \tag{2.5}$$

After some manipulation, we can see

$$\frac{p(\mathbf{X})}{1 - p(\mathbf{X})} = e^{\mathbf{X}\vec{\beta}} \tag{2.6}$$

The left-hand side of Equation 2.6 is called odds, it can take a value between 0 and inf indicating very low or very high probabilities of having label 1 respectively. After applying the logarithm to both sides of Equation 2.6, the log-odds or logit is obtained which is linear in X.

Since β_0 and all β_i are unknown, they must be estimated on the available training data. The method that is used to achieve this is called maximum likelihood estimation (MLE). MLE tries to find estimates for the coefficients β such that the estimated probability $\hat{p}(x_i)$ of having label 1 for each individual is as close as possible to the individual's observed label. In other words, MLE tries to find estimates for the coefficients β of the model $P(\mathbf{X})$ described in Equation 2.5 such that the resulting probability is close to one if the individual's observed label is 1 and close to zero if the individual's label is 0. This can be formalized as a likelihood function as shown in Equation 2.7.

$$L(\vec{\beta}|\vec{y}, \mathbf{X}) = \prod_i p(x_i)^{y_i}(1 - p(x_i))^{(1-y_i)} \tag{2.7}$$

The estimates for β are then chosen to maximize the likelihood function.[2]

2.4   Model training

The two models are trained using 5-fold nested cross validation. In each outer fold of the nested cross validation, the data is imputed using MICE and then split into training and test data. Since MICE usually generates several imputed data sets to account for the uncertainty in the imputation process, there are 5 different models for each classifier per fold. The inner fold of cross validation is used to determine the best hyperparameter combination for the classifiers using Grid Search which is also called hyperparameter tuning. The parameters for Random Forest are as follows:

- Number of trees in the forest: [10, 100, 500]

- Number of features to consider looking for best split: [4, 8, 12, 16]

- Quality of split measure: Gini impurity or Information gain

For logistic regression, Grid Search searches for an optimum among the following hyperparameters:

- Inverse of regularization strength C: [0.001, 0.01, 0.1, 1, 10, 100, 1000],

- Penalty: L1 or L2 regularization

- Algorithm to use during optimization problem: liblinear or saga

- Maximal iteration: [200, 300, 400, 500]

The remaining hyperparameters use the default values specified in the documentation of sklearn for random forests and logistic regression.

After the training, the models are then evaluated using accuracy, balanced accuracy, sensitivity, specificity, positive predictive value, and area under the ROC curve. The hyperparameters of the best performing models for Random Forest and Logistic Regression respectively are then used to obtain the final classifiers. For the validation, the whole data set is split into training and validation data according to a 80:20 split. The data set is imputed before the split due to the working of the `mice()` package in R. The classifiers are trained on the data set and then validated using the validation data. Additionally to the evaluation metrics above, the final classifiers were also evaluated using the ROC curve and the precision-recall curve. Additionally, the precision-recall curve is used to assess the performance of the classifiers regarding the slight imbalance in the data set. To further gain some insights into the classification reasoning of the random forest a decision tree is trained.

To implement this procedure, the Python libraries scikit-learn, pandas, and numpy are used. To implement the imputation method, the R library `mice()` is used and connected to the Python implementation using a bridge package called rpy2. The R implementation of the imputation method is used instead of a Python implementation since it provides more functions to inspect the MICE procedure and since it is able to handle non-normally distributed data. To visualize the results the libraries seaborn and matplotlib are used.

> Describe here k-fold nested cross validation, evaluation metrics used
>
> Here maybe also implementation of MICE in Python using rpy2

2.5 Evaluation Metrics

To describe the the evaluation metrics, the following terms will be used in the equations: True Positive, False Positive, False Negative, and True Negative. They are further characterized in Table 2.1.

		True condition	
		Condition positive	Condition negative
Predicted condition	Condition predicted positive	True Positive (TP)	False Positive (FP)
	Condition predicted negative	False Negative (FN)	True Negative (TN)

Table 2.1: Confusion matrix for binary classification

Accuracy describes the closeness of an individual's predicted label to its true label.

It is calculated as shown in Equation 2.8. Accuracy alone is not a sensible measure to determine the goodness of fit for a model since the model can be affected by imbalanced data or poor parameter initialization. Accordingly, other measures should be considered during model selection.[10]

$$accuracy = \frac{TP + TN}{TP + FP + FN + TN} \tag{2.8}$$

Balanced accuracy is used in addition to the average accuracy due to the slight imbalance in the data set. If the balanced accuracy differs significantly from the average accuracy, the imbalance has an effect on the classifier regarding class prevalence. The balanced accuracy is defined as the average of sensitivity and specificity. [10, 5]

$$balanced\ accuracy = \frac{\frac{TP}{TP+FN}\frac{TN}{TN+FP}}{2} \tag{2.9}$$

Sensitivity is interpreted as the proportion of people infected with a disease who will test positive for this disease. It is calculated as shown in Equation 2.10. Since it is only calculated using the part of the population with the disease, sensitivity can only give evidence about the true positive rate and not the false positive rate of a test.[1]

$$sensitivity = \frac{TP}{TP + FN} \tag{2.10}$$

Specificity is defined as the proportion of people without the disease that are identified correctly, i.e., who tested negative for a disease. It is calculated as show in Equation 2.11. Specificity only gives information about the proportion of people without the disease and a negative test and therefore cannot disclose anything about the false negative rate.[1]

$$specificity = \frac{TN}{TN + FP} \tag{2.11}$$

The positive predictive value (PPV) measures the post-test probability of the disease given a positive test and therefore refers to the probability that a subject with a positive test has indeed the disease. It is calculated as shown in Equation 2.12.[1]

$$PPV = \frac{TP}{TP + FP} \tag{2.12}$$

The receiver operating characteristic curve (ROC curve) is a commonly used metric to determine the performance of a model describing the tradeoff between specificity and sensitivity. It is calculated using the prediction scores of a model that are either discriminant values or posterior probabilities.[12] Most models do not produce a classification label as their out, but rather a prediction score that is then thresholded

to correspond to one of the levels of the target variable. The metrics mentioned above are all so called single-thresholded metrics since they are defined only for the thresholded output of the classifier.[12, 10] Since ROC uses the prediction scores instead of the thresholded output of a model, it can compute the sensitivity and specificity for every sensible threshold. The resulting values are then plotted with sensitivity or true positive rate as the vertical axis and the false positive rate calculated as 1 - specificity as the horizontal axis. The points are interpolated linearly to create the ROC curve. [12] Usually, a reference line is added to the plot illustrating the performance of a model that only makes random predictions, the curve for the trained model is expected to be above the reference line at all time. A model with perfect performance will appear in the very top left hand corner of the ROC space where sensitivity is 1 and false positive rate is 0.[10, 9]

The area under the ROC curve can also be measured and is called either AUC for area under the ROC curve or ROC index, it can give a numeric summary of the ROC curve. It is calculated using the integral of the ROC curve which can be easily done using the trapezoidal method since the ROC curve is discrete and stepped. A large AUC value indicates a better model performance. As a rule of thumb, an AUC over 0.7 indicates a strong predicitve model and an AUC below 0.6 indicates a weak model. AUC and ROC both are quite robust against imbalanced data, but there are other non-single-threshold metrics that can be used in this case.

The Precision-Recall curve (PRC) is another non-single-threshold metric that is especially useful when dealing with an imbalanced data set. It calculates the precision and recall (or sensitivity) values for each feasible threshold. The resulting values are then plotted with precision, also called positive predictive value, as the vertical axis and recall, also called sensitivity, as the horizontal axis, the points are interplated non-linearly to create a curve. In contrast to the ROC curve which uses a fixed reference line, the PRC's baseline is determined by the ratio of positive (P) and negative (N) samples in the data set and represented as $y = \frac{P}{P+N}$. Each point in the PRC plot corresponds to exactly one point in the ROC plot since one of the measures in the plot is the same. The area under the curve can also be calculated for the precision-recall curve similarly to the AUC of the ROC curve and is then denoted by AUC (PRC) or average precision. Similarly, a large value for the AUC (PRC) indicates a strong model and a lower value indicates a weak model.[8, 12]

add ROC, Precision-Recall-Curve

Chapter 3

Results

3.1 Results of own implementation

3.2 Comparison with original paper

Chapter 4

Discussion

4.1 Discussion of Results

4.2 Discussion of Methods

4.3 Scenarios for real-world validation

Declaration

I declare that..

Appendix A

Supplementary Tables

Feature	Data Type
Gender	Categorical
Age	Numerical (discrete)
WBC (White blood cell count)	Numerical (continuous)
Platelets	Numerical (continuous)
Neutrophils	Numerical (continuous)
Lymphocytes	Numerical (continuous)
Monocytes	Numerical (continuous)
Eosinophils	Numerical (continuous)
Basophils	Numerical (continuous)
CRP (C-reactive protein)	Numerical (continuous)
AST (aspartate aminotransferase)	Numerical (continuous)
ALT (alanine aminotransferase)	Numerical (continuous)
ALP (alkaline phosphatase)	Numerical (continuous)
GGT (gamma glutamyl transferase)	Numerical (continuous)
LDH (lactate dehydrogenase)	Numerical (continuous)
SWAB	Categorical

Table A.1: Overview over all features of the data set

Feature	Unit	Mean	Std	Median
Age	Years	61.33	18.05	64
White Blood Cell Count (WBC)	10^9/L	8.49	4.89	7.10
Platelets	10^9/L	224.91	102.61	204.00
Neutrophils	10^9/L	4.64	4.50	3.90
Lymphocytes	10^9/L	0.88	0.87	0.80
Monocytes	10^9/L	0.45	0.44	0.40
Eosinophils	10^9/L	0.04	0.12	0.00
Basophils	10^9/L	0.01	0.03	0.00
C-reactive protein (CRP)	mg/L	88.93	94.32	53.10
Aspartate Aminotransferase (AST)	U/L	53.81	57.59	36.00
Alanine Aminotransferase (ALT)	U/L	42.82	45.43	30.00
Alkaline Phosphatase (ALP)	U/L	42.21	75.71	68.00
Gamma Glutamyl Transferase (GGT)	U/L	40.20	101.29	0.00
Lactate dehydrogenase (LDH)	U/L	264.54	238.53	254.00

Table A.2: Descriptive statistics for numerical features in data set (including missing values as in [5])

Feature	Number of NaN (in %)
Gender	0 (0 %)
Age	2 (0.72 %)
WBC (White blood cell count)	2 (0.72 %)
Platelets	2 (0.72 %)
Neutrophils	70 (25.09 %)
Lymphocytes	71 (25.45 %)
Monocytes	70 (25.09 %)
Eosinophils	70 (25.09 %)
Basophils	71 (25.45 %)
CRP (C-reactive protein)	6 (2.15 %)
AST (aspartate aminotransferase)	2 (0.72 %)
ALT (alanine aminotransferase)	13 (4.66 %)
ALP (alkaline phosphatase)	148 (53.05 %)
GGT (gamma glutamyl transferase)	143 (51.25 %)
LDH (lactate dehydrogenase)	85 (30.47 %)
SWAB	0 (0 %)

Table A.3: Number of missing values and their proportion the total number of data points

Feature	Unit	Mean	Std	Median
Age	Years	61.78	17.81	64
White Blood Cell Count (WBC)	10^9/L	8.55	4.86	7.10
Platelets	10^9/L	226.5	101.2	205.00
Neutrophils	10^9/L	6.20	4.17	5.10
Lymphocytes	10^9/L	1.19	0.80	1.00
Monocytes	10^9/L	0.61	0.41	0.50
Eosinophils	10^9/L	0.06	0.13	0.00
Basophils	10^9/L	0.01	0.04	0.00
C-reactive protein (CRP)	mg/L	90.89	94.42	54.20
Aspartate Aminotransferase (AST)	U/L	54.20	57.61	36.00
Alanine Aminotransferase (ALT)	U/L	44.92	45.50	31.00
Alkaline Phosphatase (ALP)	U/L	89.89	89.09	71.00
Gamma Glutamyl Transferase (GGT)	U/L	82.48	132.70	41.00
Lactate dehydrogenase (LDH)	U/L	380.45	193.98	328.00

Table A.4: Descriptive statistics for numerical features in data set (excluding missing values)

Feature	Test statistic	p-value
Age	0.976	0.000
WBC (White blood cell count)	0.873	0.000
Platelets	0.930	0.000
Neutrophils	0.838	0.000
Lymphocytes	0.785	0.000
Monocytes	0.811	0.000
Eosinophils	0.457	0.000
Basophils	0.395	0.000
CRP (C-reactive protein)	0.836	0.000
AST (aspartate aminotransferase)	0.556	0.000
ALT (alanine aminotransferase)	0.629	0.000
ALP (alkaline phosphatase)	0.420	0.000
GGT (gamma glutamyl transferase)	0.500	0.000
LDH (lactate dehydrogenase)	0.877	0.000

Table A.5: Results of Shapiro-Wilk test for normality for numerical features of the data set

Appendix B

Supplementary Figures

Figure B.1: Kernel density plots of the numerical features of the data set

Nested k-fold Cross-Validation

1. Define set of hyper-parameter combinations, **C**, for current model. If model has no hyper-parameters, **C** is the empty set.
2. Divide data into K folds with approximately equal distribution of cases and controls
3. (**outer loop**) For fold k_i in the K folds:
   1. Set fold k_i as the test set
   2. Perform automated feature selection on the remaining $K-1$ folds
   3. For parameter combination c in **C**:
      1. (**inner loop**) For fold k_j in the remaining $K-1$ folds:
         1. Set fold k_j as the validation set
         2. Train model on remaining $K-2$ folds
         3. Evaluate model performance on fold k_j
      2. Calculate average performance over $K-2$ folds for parameter combination c
   4. Train model on $K-1$ folds using hyper-parameter combination that yielded best average performance over all steps of the **inner loop**
   5. Evaluate model performance on fold k_i.
4. Calculate average performance over K folds

Figure B.2: K-fold nested cross validation

Appendix C

Supplementary Formulas

Bibliography

[1] A. K. Akobeng. "Understanding diagnostic tests 1: sensitivity, specificity and predictive values". In: *Acta Paediatr* 96.3 (2007), pp. 338–41. ISSN: 0803-5253 (Print) 0803-5253 (Linking). DOI: 10.1111/j.1651-2227.2006.00180.x. URL: https://www.ncbi.nlm.nih.gov/pubmed/17407452.

[2] *An Introduction to Statistical Learning*. Springer Texts in Statistics. 2013. ISBN: 978-1-4614-7137-0 978-1-4614-7138-7. DOI: 10.1007/978-1-4614-7138-7.

[3] M. J. Azur et al. "Multiple imputation by chained equations: what is it and how does it work?" In: *Int J Methods Psychiatr Res* 20.1 (2011), pp. 40–9. ISSN: 1557-0657 (Electronic) 1049-8931 (Linking). DOI: 10.1002/mpr.329. URL: https://www.ncbi.nlm.nih.gov/pubmed/21499542,%20https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3074241/pdf/MPR-20-40.pdf.

[4] Leo Breiman. "Random Forests". In: *Machine Learning* 45.1 (2001), pp. 5–32. ISSN: 08856125. DOI: 10.1023/a:1010933404324. URL: https://link.springer.com/content/pdf/10.1023/A:1010933404324.pdf.

[5] D. Brinati et al. "Detection of COVID-19 Infection from Routine Blood Exams with Machine Learning: A Feasibility Study". In: *J Med Syst* 44.8 (2020), p. 135. ISSN: 1573-689X (Electronic) 0148-5598 (Linking). DOI: 10.1007/s10916-020-01597-4. URL: https://www.ncbi.nlm.nih.gov/pubmed/32607737,%20https://www.ncbi.nlm.nih.gov/pmc/articles/PMC7326624/pdf/10916_2020_Article_1597%20.pdf,https://zenodo.org/record/3886927#.X3xhPO1CRPY.

[6] Stef van Buuren. *Flexible Imputation of Missing Data, Second Edition*. 2nd ed. Boca Raton, Fla: CRC Press, 2018. ISBN: 978-0-429-96034-5. URL: https://stefvanbuuren.name/fimd/.

[7] Stef van Buuren and Karin Groothuis-Oudshoorn. "mice: Multivariate Imputation by Chained Equations inR". In: *Journal of Statistical Software* 45.3 (2011). ISSN: 1548-7660. DOI: 10.18637/jss.v045.i03.

[8]   Jesse Davis and Mark Goadrich. "The relationship between Precision-Recall and ROC curves". In: *Journal* (2006), pp. 233–240. ISSN: ISSN. DOI: `10.1145/1143844.1143874`. URL: `https://dl.acm.org/doi/pdf/10.1145/1143844.1143874`.

[9]   Tom Fawcett. "An introduction to ROC analysis". In: *Pattern Recognition Letters* 27.8 (2006), pp. 861–874. ISSN: 01678655. DOI: `10.1016/j.patrec.2005.10.010`. URL: `https://www.sciencedirect.com/science/article/abs/pii/S016786550500303X?via%3Dih%20ub`.

[10]   John D. Kelleher, Aoife DArcy, and Brian Mac Namee. *Fundamentals of Machine Learning for Predictive Data Analytics: algorithms, worked examples, and case studies.* 1st ed. Camebridge, MA: The MIT Press, 2015. URL: `https://mitpress.mit.edu/books/fundamentals-machine-learning-predictive-data-ana%20lytics`.

[11]   K. Kühnberger and N. Potyka. *Methods of Artifical Intelligence.* Lecture. 2018.

[12]   T. Saito and M. Rehmsmeier. "The precision-recall plot is more informative than the ROC plot when evaluating binary classifiers on imbalanced datasets". In: *PLoS One* 10.3 (2015), e0118432. ISSN: 1932-6203 (Electronic) 1932-6203 (Linking). DOI: `10.1371/journal.pone.0118432`. URL: `https://www.ncbi.nlm.nih.gov/pubmed/25738806,%20https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4349800/pdf/pone.0118432.pdf`.

[13]   I. R. White, P. Royston, and A. M. Wood. "Multiple imputation using chained equations: Issues and guidance for practice". In: *Stat Med* 30.4 (2011), pp. 377–99. ISSN: 1097-0258 (Electronic) 0277-6715 (Linking). DOI: `10.1002/sim.4067`. URL: `https://www.ncbi.nlm.nih.gov/pubmed/21225900,%20https://onlinelibrary.wiley.com/doi/abs/10.1002/sim.4067`.